



# Data Service Server Monitoring with Prometheus

## User's Manual

V1.0.2

## Revision History

Date	Version	Author	Reviewer	Description
2019-03-05	1.0.0	Evelyn Tang Evelyn.Tang@advantech.com.tw		First version released
2019-03-07	1.0.1	Evelyn Tang Evelyn.Tang@advantech.com.tw		Paragraph adjustment Add architecture image
2019-03-08	1.0.2	Evelyn Tang Evelyn.Tang@advantech.com.tw		Change name to "Data Service Server"

# Table of Contents

1	General Introduction .....	4
2	Prometheus Overview.....	6
3	Data visualization.....	8
	FAQ.....	20
	Reference.....	23

Confidential

# Table of Figures

Figure 1 Architecture of Data Service Server .....	4
Figure 2 Architecture of Prometheus monitoring system [2] .....	6
Figure 3 Grafana login page .....	8
Figure 4 Dashboard overview .....	10
Figure 5 Kubernetes Cluster Overview .....	11
Figure 6 Node Exporter Overview .....	12
Figure 7 Kubernetes Pods .....	13
Figure 8 RabbitMQ Overview .....	15
Figure 9 PostgreSQL Overview .....	16
Figure 10 Prometheus Web UI.....	17
Figure 11 Simple query result.....	18
Figure 12 Advanced query result.....	19

# 1 General Introduction

## 1.1 Data Service Server Architecture



Figure 1 Architecture of Data Service Server

Data Service Server runs Kubernetes for automating deployment, scaling and management of containerized applications. [1] As illustrated in Figure 1, collecting data from an environment composed of so many moving parts is complex and Prometheus is the best monitoring and alerting tool for Kubernetes and Docker.

Prometheus was built specifically to monitor applications and microservices running in containers at scale and is native to containerized environments. Originally developed at Soundcloud, pioneer in the adoption of cloud technology. [2] In 2016,

Prometheus project became the second hosted project of the Cloud Native Computing Foundation (CNCF) after Kubernetes.

## 1.2 Prometheus and Kubernetes



Prometheus is a pull-based monitoring system, which means that central Prometheus servers discover and pull metrics from your services. The discovery and pull system fits well with a dynamic, cloud native environment such as Kubernetes, where Prometheus integrates well with Kubernetes to discover and enumerate the services you have running. As you scale up a service, Prometheus automatically starts pulling metrics from the extra replicas. Similarly as nodes fail and pods are restarted, Prometheus automatically notices and scrapes them. [3] The architecture of Prometheus will be briefly introduced in next section.



## 2 Prometheus Overview

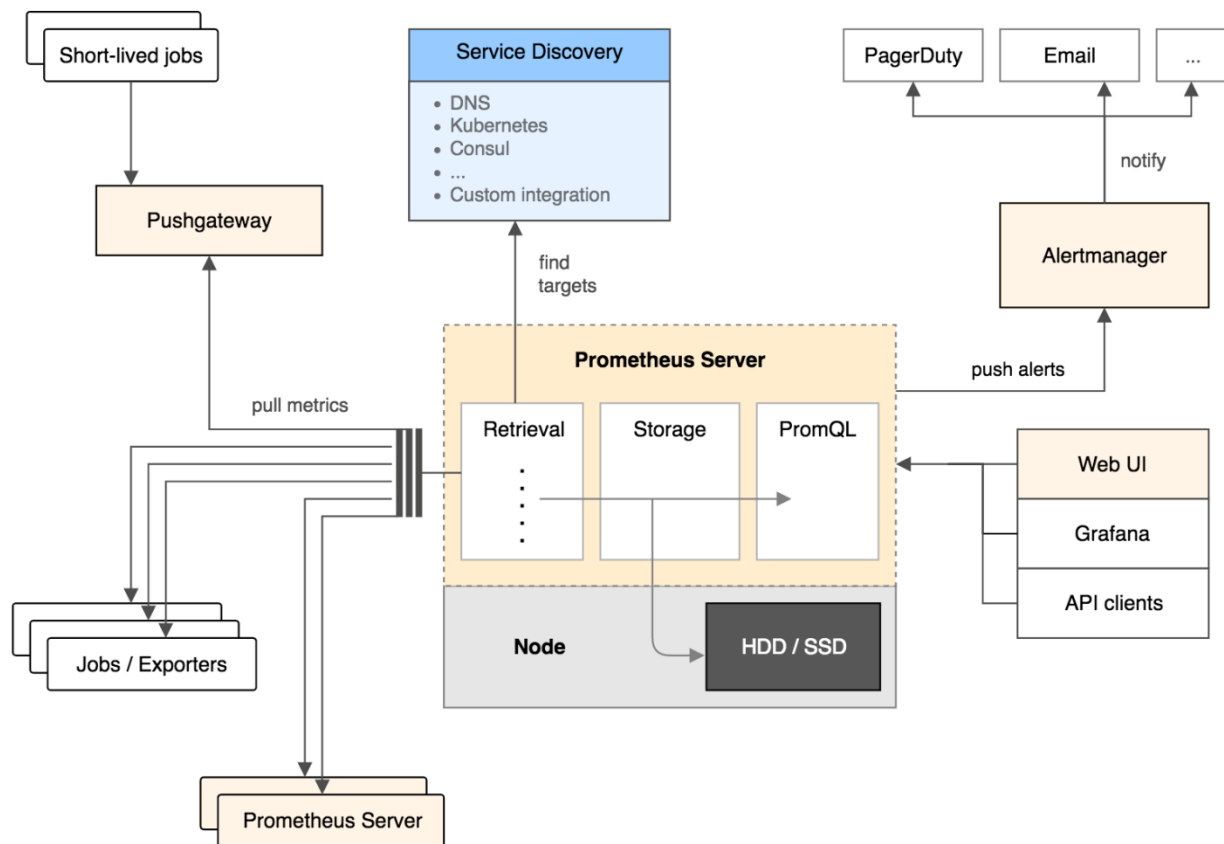


Figure 2 Architecture of Prometheus monitoring system [4]

As shown in Figure 2, Prometheus consists of several components working together. Prometheus server pulls in metrics at regular intervals through HTTP endpoints which provide information about hosts and various applications. These endpoints may be provided by the monitored application itself, or through an “exporter” . The endpoint URL is usually `/metrics`. [5]

Since short-lived jobs not exist long enough to be scraped, they can instead push their metrics to a Pushgateway. Pushgateway then exposes these metrics to

Prometheus server. [6] Last but not least, Prometheus Server also exposes its own metrics and monitors its own metrics. [7]

After Prometheus server scrapes metrics from instrumented jobs, either directly or via an intermediary push gateway for short-lived jobs, it stores time-series data locally. User can use PromQL to query this data or send alerts to the Alertmanager, which will convert them into pages, emails, and other notifications. [8]

Prometheus provides a web interface (Web UI) to run queries. Other applications also can run queries through the HTTP API to retrieve and work with the data; we apply Grafana to visualize the data.

Next section will cover two parts: Grafana dashboards and Prometheus WebUI. It is recommended that users completely read this document before evaluation since you will understand which data visualization method meets your needs.



## 3 Data visualization

### 3.1 Grafana Dashboards

#### 3.1.1 Introduction

Grafana is the leading graph and dashboard builder for visualizing time series infrastructure and application metrics. Therefore, we apply Grafana as data visualization tool.

#### 3.1.2 Imported Dashboards

First of all, go to your Grafana main view and login admin account.

We take <http://portal.grafana.example.com> for example as Figure 3.

Please go to your Grafana domain server.

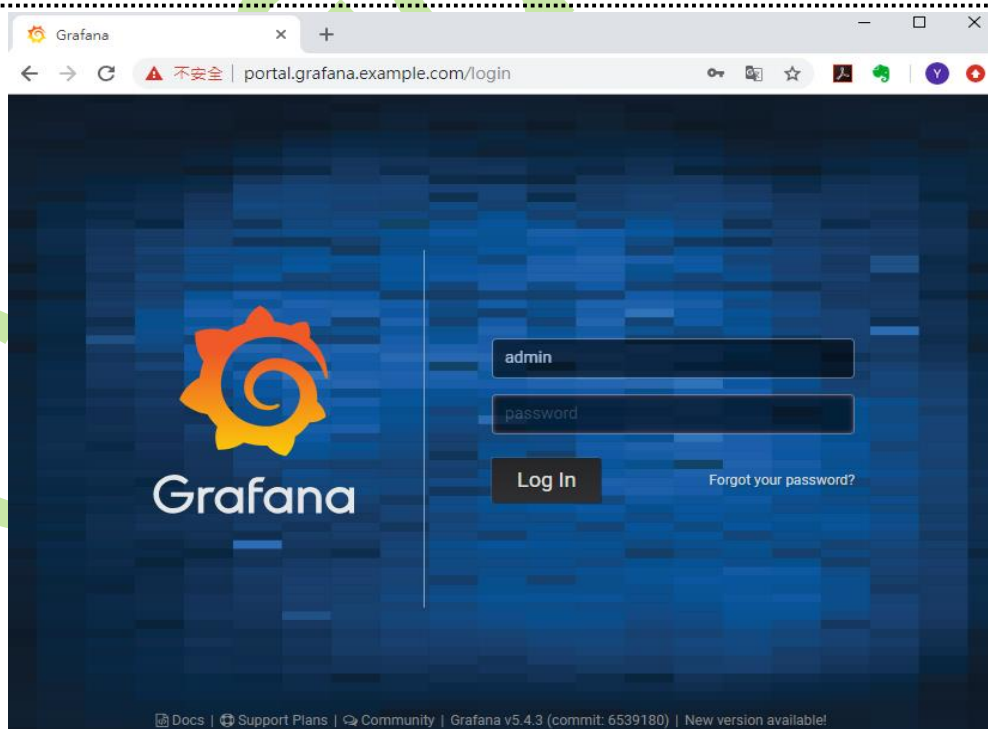


Figure 3 Grafana login page

Please access admin account as below.

User Name: **admin**

Password: **P@ssw0rd**

You can get username and password of admin account by issuing the following **kubectl** command.

✧ Get Username:

```
# kubectl get secret --namespace monitoring grafana -o  
jsonpath="{.data.admin-user}" | base64 --decode ; echo
```

✧ Get Password:

```
# kubectl get secret --namespace monitoring grafana -o  
jsonpath="{.data.admin-password}" | base64 --decode ; echo
```

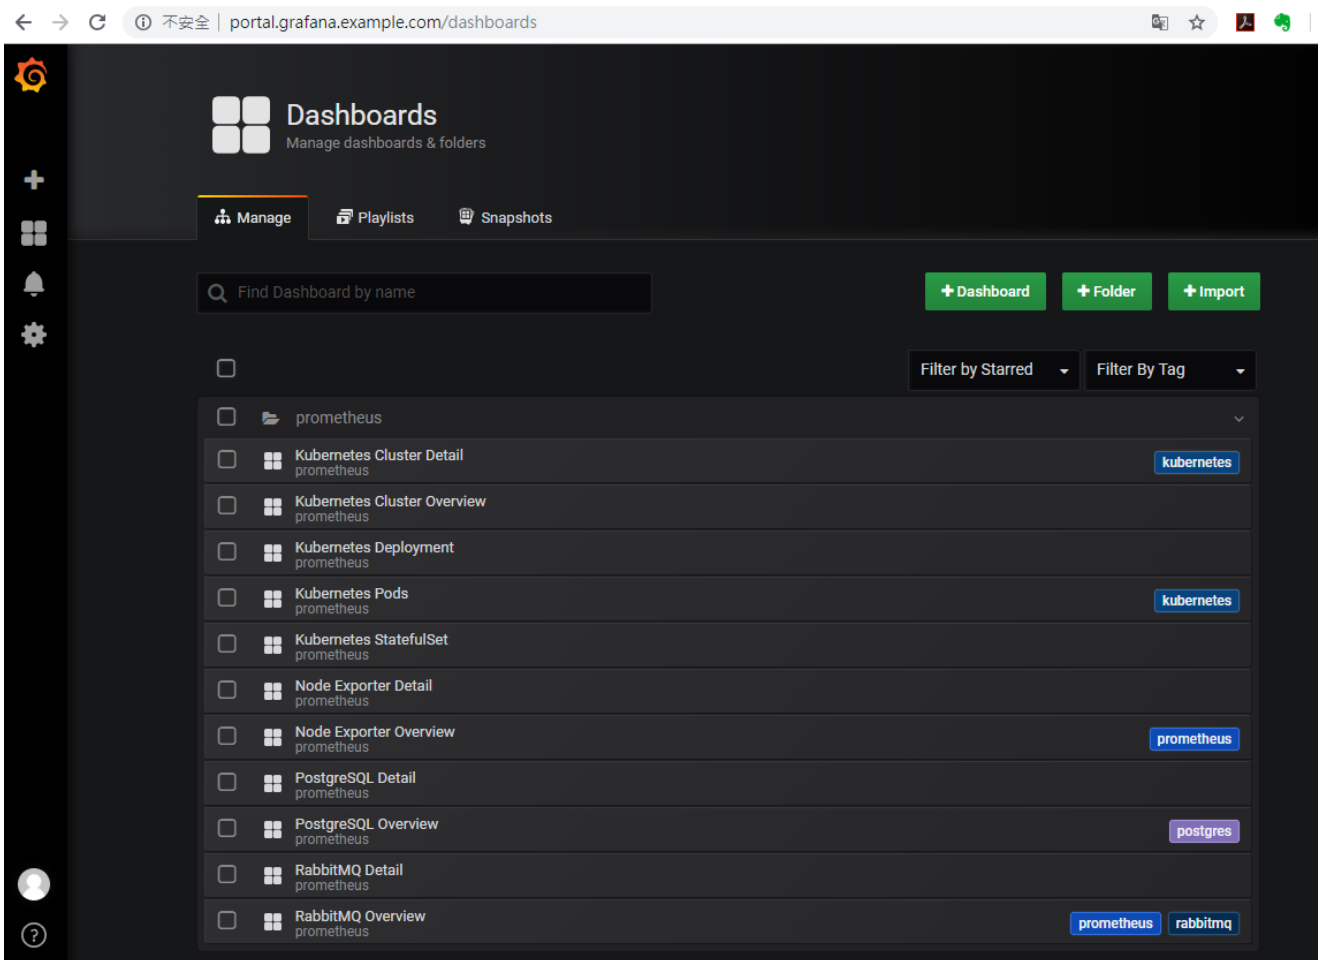


Figure 4 Dashboard overview

Generally speaking, there are several Kubernetes metrics worthy to monitor. As shown in Figure 4, adequate and detailed dashboards have been imported.

- Kubernetes Cluster Overview

For cluster monitoring, the objective is to monitor the health of the entire Kubernetes cluster. As an administrator, we are interested in discovering if all the nodes in the cluster are working properly and at what capacity, how many applications are running on each node, and the resource utilization of the entire cluster. [9]The overview of cluster is shown in Figure 5 below.



Figure 5 Kubernetes Cluster Overview

There are many measurable metrics worthy of mention, all related to node resource utilization, network bandwidth, disk utilization, CPU, and memory utilization are examples of this. You can navigate to dashboard “Kubernetes Cluster Detail” for more cluster information in details.

- Node Exporter Overview

One of the most widely used exporters is the NodeExporter. When NodeExporter runs on a host, it will collect data on I/O, memory, disk and CPU pressure and expose them for scraping. The overview of NodeExporter is as below Figure 6. [10] Also, we have another dashboard “Node Exporter Detail” , it shows individual details of memory, network and other system metrics.

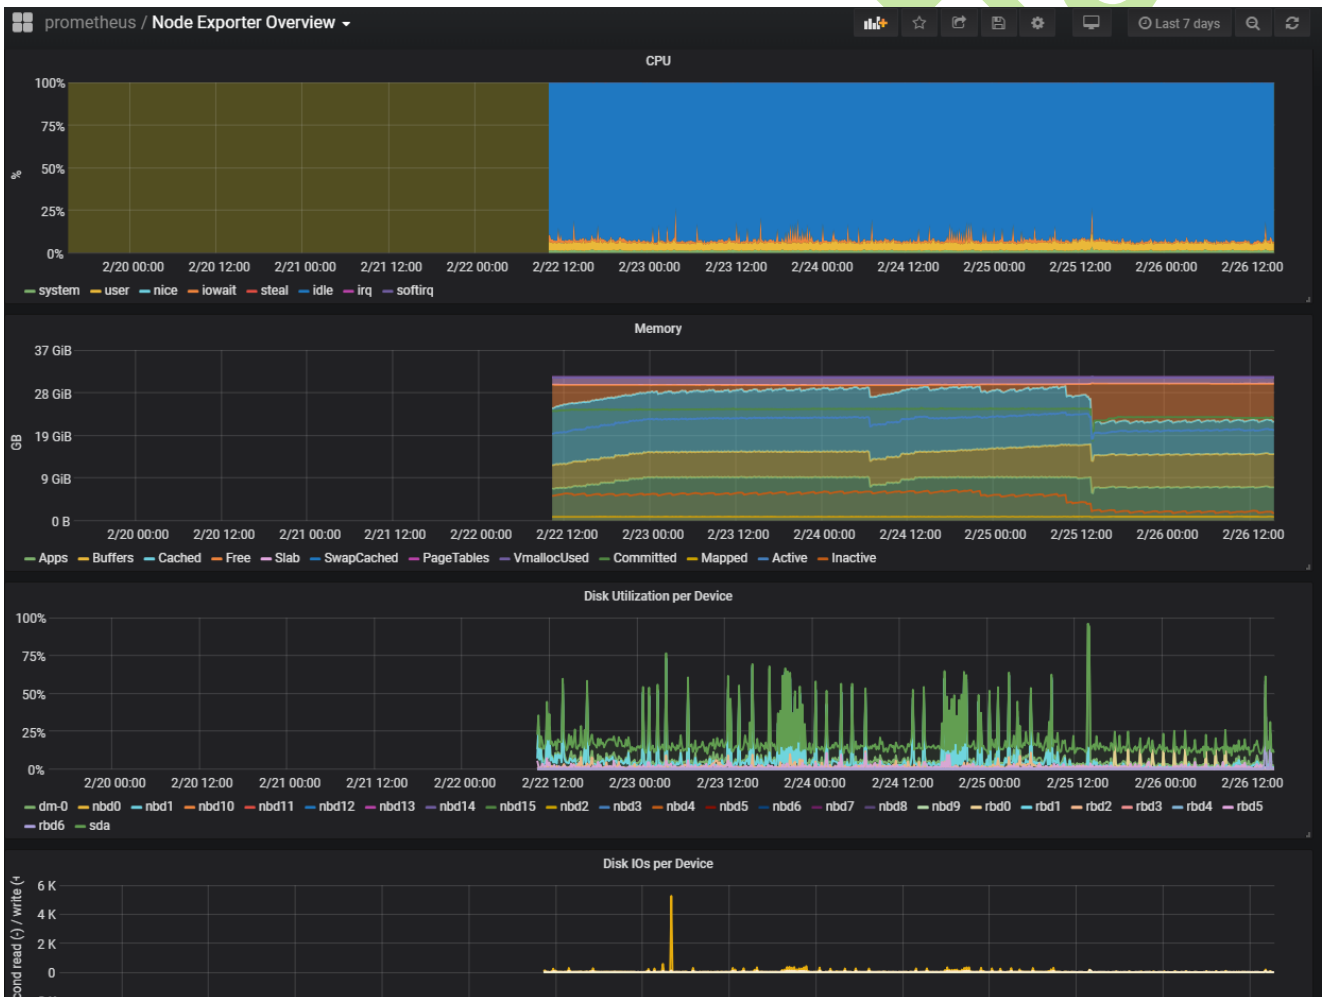


Figure 6 Node Exporter Overview

- Kubernetes Pods




Figure 7 Kubernetes Pods

The act of monitoring a pod can also extend to its containers. You can monitor individual container if there is more than one container in a pod. Container metrics query running container information like CPU, network, and memory usage. As Figure 7 shows above, maximum limitation is set to compare with memory usage.

- Other Kubernetes dashboards

**"Kubernetes Deployment"** : Kubernetes deployment is an abstraction layer for the pods. The main purpose of the deployment is to maintain the resources declared in the deployment configuration to be its desired state. [11] User can monitor the status of every deployment in every namespace.

**"Kubernetes StatefulSet"** : Manages the deployment and scaling of a set of Pods, and provides guarantees about the ordering and uniqueness of these Pods. [12] The status overview of StatefulSet will show in this dashboard.

- Third-Party Exporter

Some other services are not natively integrated, but can be easily adapted using an exporter. An exporter is a service that collects service status and translates to Prometheus metrics ready to be scraped. [13] We have two examples in the Data Service Server.

- RabbitMQ Overview



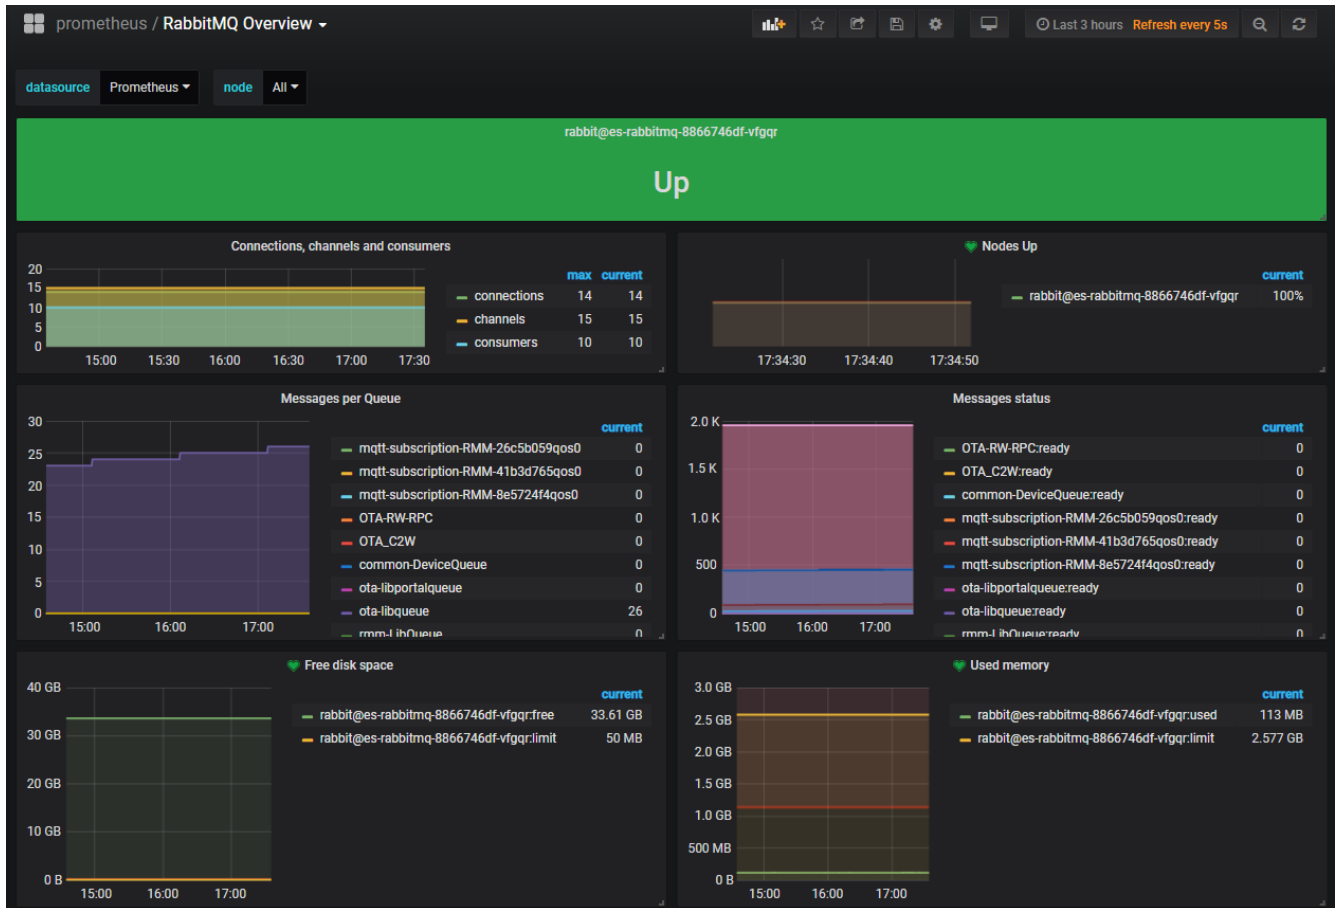


Figure 8 RabbitMQ Overview

RabbitMQ is the most widely deployed open source message broker, it supports multiple messaging protocols, message queuing, delivery acknowledgment, flexible routing to queues, multiple exchange type. [14] RabbitMQ exporter is applied to provide a starting point for monitoring RabbitMQ metrics as shown in Figure 8. Prometheus is configured 60 seconds interval to collect RabbitMQ status.

## ➤ PostgreSQL Overview

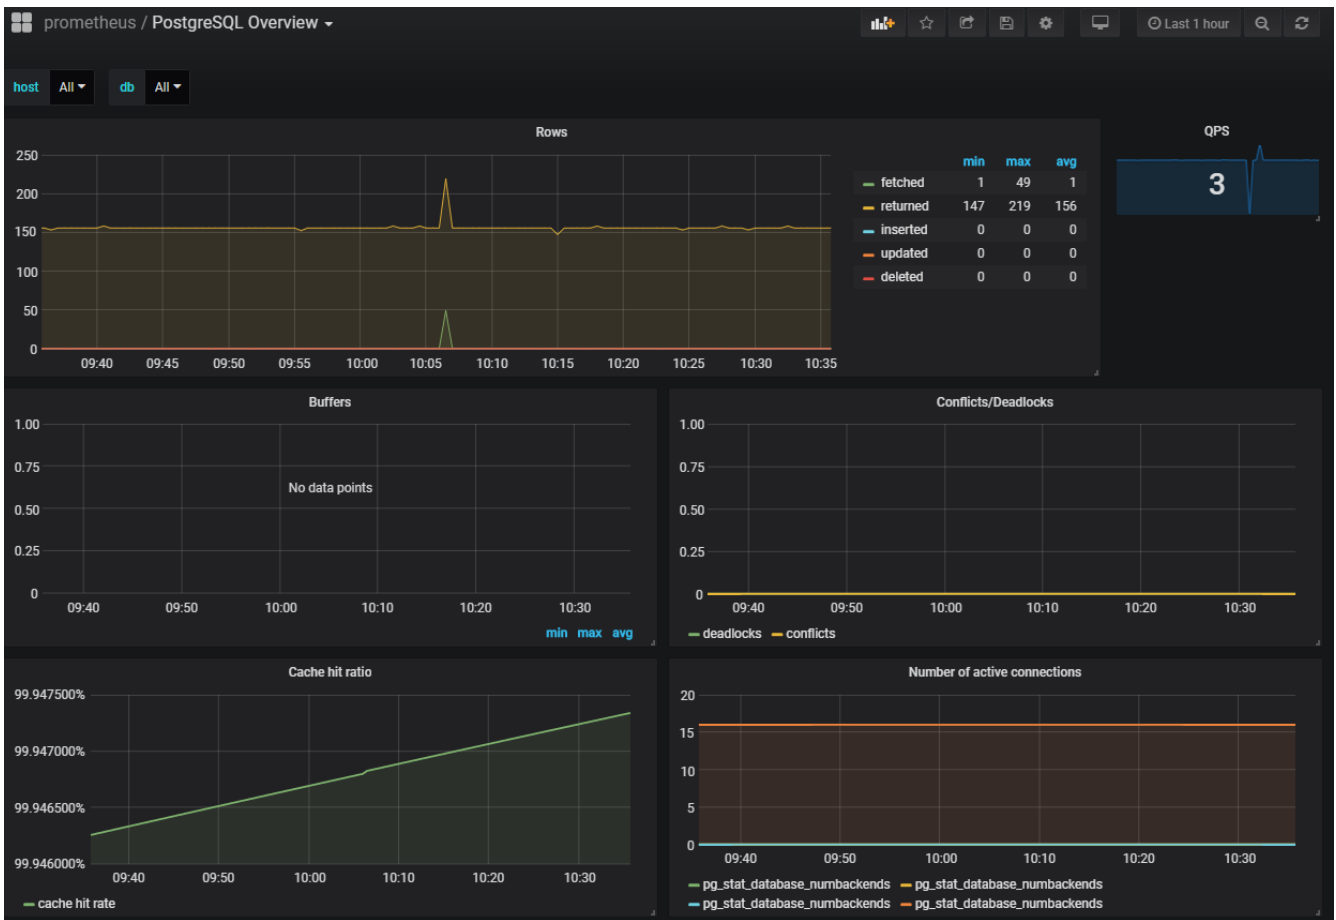


Figure 9 PostgreSQL Overview

PostgreSQL comes with many features aimed to help developers build applications, administrators to protect data integrity and build fault-tolerant environments, and help you manage your data no matter how big or small the dataset. [15] [16] The account and device information of Data Service Server are stored in PostgreSQL, as shown in Figure 9, you can monitor real-time status of PostgreSQL.

## 3.2 Prometheus Web UI

### 3.2.1 Introduction

Prometheus provides a functional query language called PromQL (Prometheus Query Language) that lets user select and aggregate time series data in real time. The result of an expression can either be shown as a graph, viewed as tabular data in Prometheus's expression browser. It will be a good choice if you need to query specific metrics using PromQL.

### 3.2.2 PromQL

At first, go to your Prometheus Web UI with no username and password.

We take <http://portal.prometheus.example.com> for example as Figure 10.

Please go to your Prometheus WebUI domain server.




Figure 10 Prometheus Web UI

Take `container_memory_usage_bytes` for quick example, it selects all time-series that have the `container_memory_usage_bytes` metric name. Figure 11 below shows the query results with graph mode.

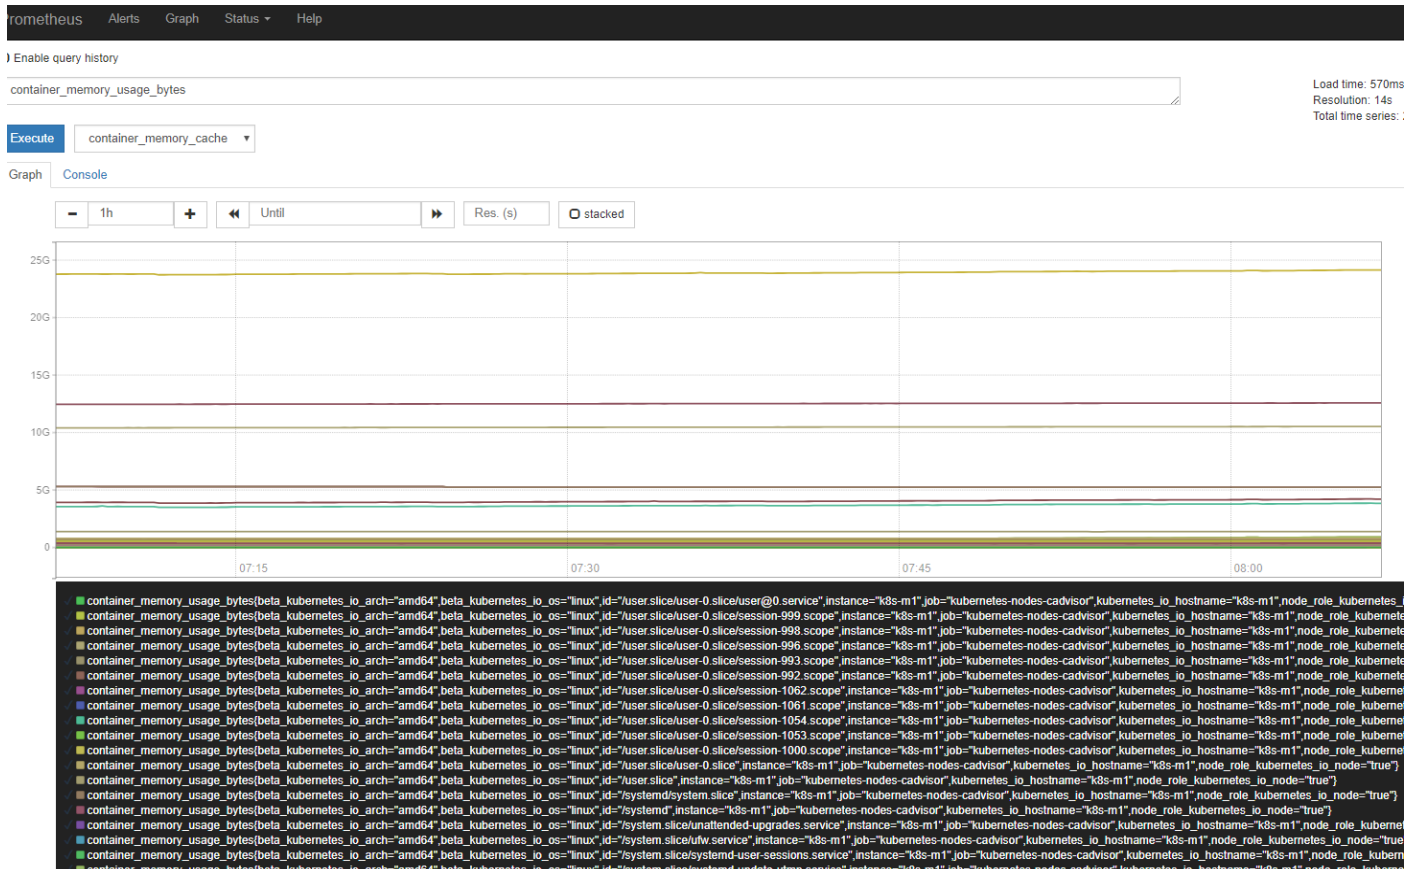
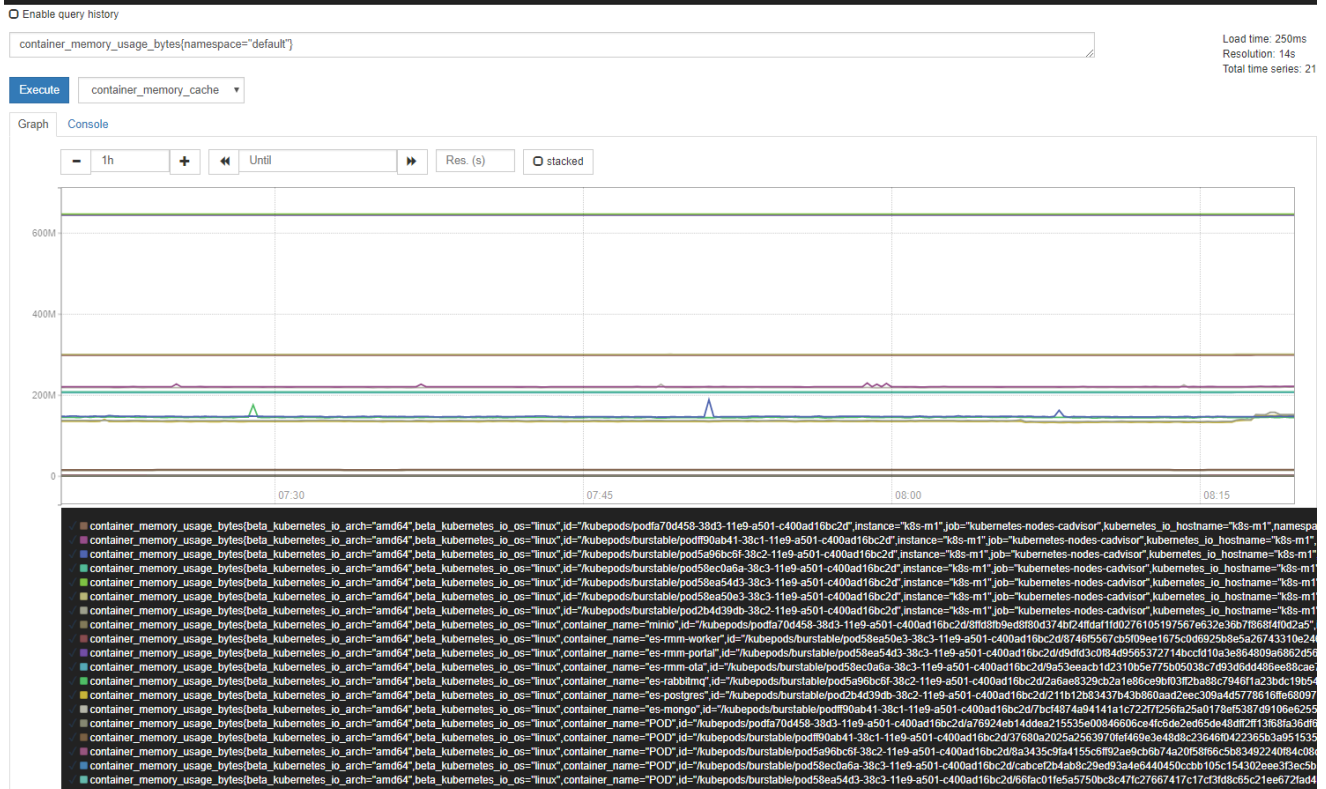


Figure 11 Simple query result

The above results contain all the `container_memory_usage_bytes` metrics, the data is large. If you want to query more specific data in advanced, append a set of labels to match in curly braces (`{}`). Take `container_memory_usage_bytes` for example, append one term in PromQL as below.

```
✧ container_memory_usage_bytes{namespace="default"}
```



### Figure 12 Advanced query result

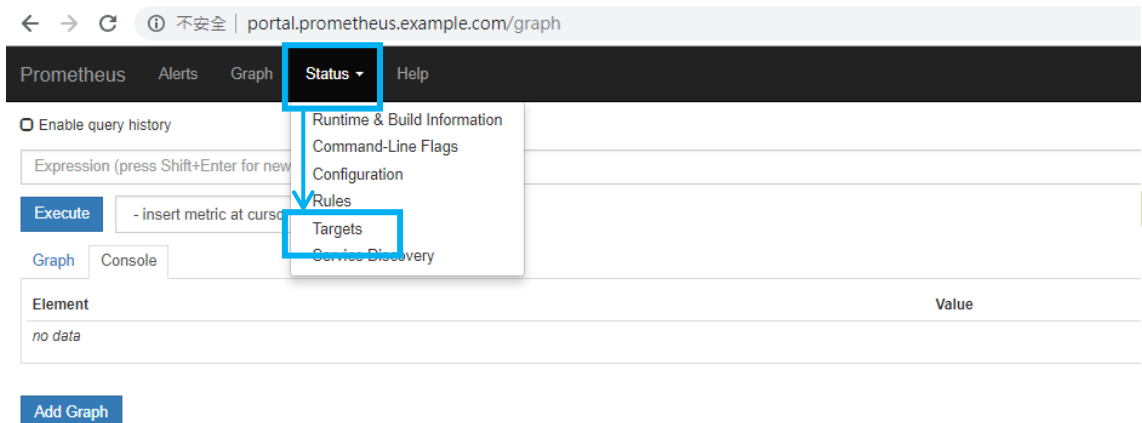
The result shows in Figure 12 that `container_memory_usage_bytes` metrics with namespace `"default"`. Therefore, user can query any specific metric if in need. Please refer to the document for more detail of PromQL.

<https://prometheus.io/docs/prometheus/latest/querying/basics/>

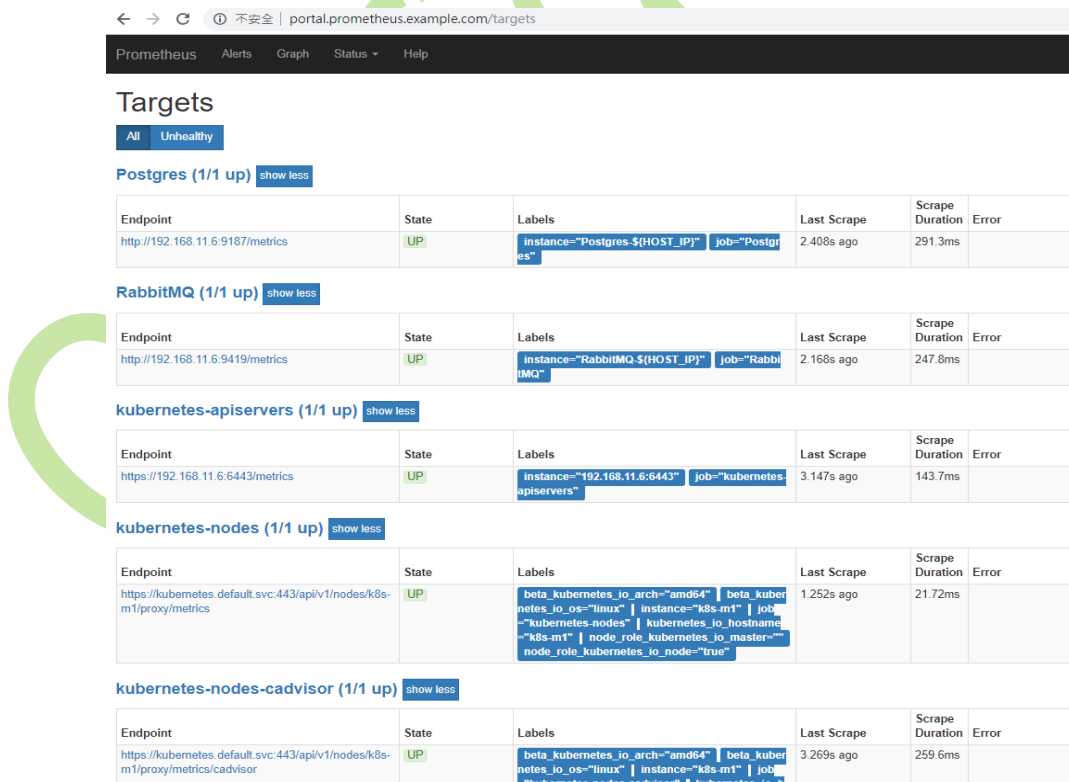
## FAQ

✧ How to simply check if all the targets are under monitored?

1. Go to your Prometheus Web UI and navigate to “Targets” page.



2. All the monitored targets will be listed as below, the targets will show “up” if their status are healthy.



✧ As mentioned above, if there is a target with “down” status and error message,

how to check the error in detail?

Please use command `kubectl` to get more information.

1. Issue command to get all resources in namespace "monitoring" .

```
# kubectl get all -n=monitoring
```

All the resources in "monitoring" will show as below.

```
root@k8s-m1:~# kubectl get all -n=monitoring
NAME                                     READY   STATUS    RESTARTS   AGE
pod/grafana-884c85f54-z7gtg             1/1     Running   1          11d
pod/postgres-exporter-prometheus-postgres-exporter-d8cf5bfb-mnrq5 1/1     Running   0          11d
pod/prometheus-alertmanager-64f94bf454-gffvm 2/2     Running   0          5d
pod/prometheus-kube-state-metrics-6d6ff7456-z8kxn 1/1     Running   0          5d
pod/prometheus-node-exporter-frhnc       1/1     Running   0          5d
pod/prometheus-pushgateway-577cd4d4d6-vc2wr 1/1     Running   0          5d
pod/prometheus-server-79d9d4b8f4-7qkl6    2/2     Running   0          5d
pod/rabbitmq-exporter-prometheus-rabbitmq-exporter-86458495dc-k5h9k 1/1     Running   0          11d

NAME                                     TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
service/grafana                         ClusterIP      10.233.41.206   <none>           3000/TCP         11d
service/postgres-exporter-prometheus-postgres-exporter ClusterIP      10.233.0.13    <none>           9187/TCP         11d
service/prometheus-alertmanager         ClusterIP      10.233.32.174   <none>           80/TCP           5d
service/prometheus-kube-state-metrics    ClusterIP      None            <none>           80/TCP           5d
service/prometheus-node-exporter         ClusterIP      None            <none>           9100/TCP         5d
service/prometheus-pushgateway           ClusterIP      10.233.38.104   <none>           9091/TCP         5d
service/prometheus-server                ClusterIP      10.233.53.84    <none>           9090/TCP         5d
service/rabbitmq-exporter-prometheus-rabbitmq-exporter ClusterIP      10.233.36.4     <none>           9419/TCP         11d
```

Take "Postgres" for example and assume "Postgres" is down.

## Targets

All Unhealthy

Postgres (1/1 down) [show less](#)

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://192.168.11.6:9187/metrics	down	instance="Postgres-HOST_IP" job="Postgres"	34.117s ago	437.6ms	

RabbitMQ (1/1 up) [show less](#)

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://192.168.11.6:9419/metrics	UP	instance="RabbitMQ-HOST_IP" job="RabbitMQ"	47.52s ago	56.58ms	

kubernetes-apiservers (1/1 up) [show less](#)

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
https://192.168.11.6:6443/metrics	UP	instance="192.168.11.6:6443" job="kubernetes-apiservers"	5.041s ago	115.1ms	

kubernetes-nodes (1/1 up) [show less](#)

2. We need to get more information from "Postgres" Pod.



```
root@k8s-m1:~# kubectl get all -n=monitoring
```

NAME	READY	STATUS	RESTARTS	AGE
pod/grafana-884c85f54-z7gtg	1/1	Running	1	11d
pod/postgres-exporter-prometheus-postgres-exporter-d8cf5bfbb-mnrq5	1/1	Running	0	11d
pod/prometheus-alertmanager-64f94bf454-qffvm	2/2	Running	0	5d
pod/prometheus-kube-state-metrics-6d6ff7456-z8kxn	1/1	Running	0	5d
pod/prometheus-node-exporter-frhnc	1/1	Running	0	5d
pod/prometheus-pushgateway-577cd4d4d6-vc2wr	1/1	Running	0	5d
pod/prometheus-server-79d9d4b8f4-7qkl6	2/2	Running	0	5d
pod/rabbitmq-exporter-prometheus-rabbitmq-exporter-86458495dc-k5h9k	1/1	Running	0	11d

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/grafana	ClusterIP	10.233.41.206	<none>	3000/TCP	11d
service/postgres-exporter-prometheus-postgres-exporter	ClusterIP	10.233.0.13	<none>	9187/TCP	11d
service/prometheus-alertmanager	ClusterIP	10.233.32.174	<none>	80/TCP	5d
service/prometheus-kube-state-metrics	ClusterIP	None	<none>	80/TCP	5d
service/prometheus-node-exporter	ClusterIP	None	<none>	9100/TCP	5d
service/prometheus-pushgateway	ClusterIP	10.233.38.104	<none>	9091/TCP	5d
service/prometheus-server	ClusterIP	10.233.53.84	<none>	9090/TCP	5d
service/rabbitmq-exporter-prometheus-rabbitmq-exporter	ClusterIP	10.233.36.4	<none>	9419/TCP	11d

```
# kubectl logs -n=monitoring ${POD_NAME}
```

Please fill in your pod name here and you will see the logs of this pod.

i.e. `kubectl logs -n=monitoring postgres-exporter-prometheus-postgres-exporter-d8cf5bfbb-mnrq5`

```
root@k8s-m1:~# kubectl logs -n=monitoring postgres-exporter-prometheus-postgres-exporter-d8cf5bfbb-mnrq5
time="2019-02-22T02:24:03Z" level=info msg="Established new database connection." source="postgres_exporter.go:995"
time="2019-02-22T02:24:03Z" level=info msg="Semantic Version Changed: 0.0.0 -> 0.6.11" source="postgres_exporter.go:925"
time="2019-02-22T02:24:03Z" level=info msg="Starting Server: :9187" source="postgres_exporter.go:1137"
time="2019-02-25T05:47:15Z" level=warning msg="Proceeding with outdated query maps, as the Postgres version could not be determined: Error scanning version string: dial tcp 192.168.11.6:5432: conn
: connection refused" source="postgres_exporter.go:1041"
time="2019-02-25T05:47:15Z" level=info msg="Error retrieving settings: Error running query on database: pg dial tcp 192.168.11.6:5432: connect: connection refused\n" source="postgres_exporter.go:
9"
time="2019-02-25T05:47:15Z" level=info msg="Error running query on database: pg_stat_database_conflicts dial tcp 192.168.11.6:5432: connect: connection refused\n" source="postgres_exporter.go:893"
```

3. By the two steps above, if you can correct the error successfully, congrats!

But if you still have no ideas, please issue command below.

```
# kubectl describe pod -n=monitoring ${POD_NAME}
```

Please fill in your pod name here.

i.e. `kubectl describe pod -n=monitoring postgres-exporter-prometheus-postgres-exporter-d8cf5bfbb-mnrq5`

Please copy and paste all the information from “kubectl logs” and “kubectl describe”

commands and contact Advantech support team. Thanks!

## Reference

- [1] "Kubernetes," [Online]. Available: <https://kubernetes.io/>.
- [2] Weaveworks, "Monitoring Kubernetes with Prometheus," [Online]. Available: <https://www.weave.works/technologies/monitoring-kubernetes-with-prometheus/>.
- [3] Tom, "Prometheus and Kubernetes: A Perfect Match," [Online]. Available: <https://www.weave.works/blog/prometheus-kubernetes-perfect-match/>.
- [4] 殷. 蔡. 吴莉, "Prometheus 入门与实践," [Online]. Available: <https://www.ibm.com/developerworks/cn/cloud/library/cl-lo-prometheus-getting-started-and-practice/index.html>.
- [5] S. BISWAS. [Online]. Available: <https://www.booleanworld.com/install-use-prometheus-monitoring/>.
- [6] "Github of Prometheus/pushgateway," [Online]. Available: <https://github.com/prometheus/pushgateway>.
- [7] K. Aziz. [Online]. Available: <http://weblogs.com.pk/khurram/archive/2018/04/16/prometheus.aspx>.
- [8] B. Brazil, Prometheus: Up & Running, O'Reilly Media, Inc., 2018.
- [9] D. Berman, "Kubernetes Monitoring: Best Practices, Methods, and Existing Solutions," [Online]. Available: <https://logz.io/blog/kubernetes-monitoring/>.
- [10] A. Ellis, "Monitor your applications with Prometheus," [Online]. Available: <https://blog.alexellis.io/prometheus-monitoring/>.
- [11] devopscube, "Kubernetes Deployment Tutorial For Beginners," [Online]. Available: <https://devopscube.com/kubernetes-deployment-tutorial/>.
- [12] "Kubernetes concepts: StatefulSets," [Online]. Available: <https://kubernetes.io/docs/concepts/workloads/controllers/statefulset/>.
- [13] M. Burillo, "Kubernetes Monitoring with Prometheus -The ultimate guide," [Online]. Available: <https://sysdig.com/blog/kubernetes-monitoring-prometheus/>.
- [14] Kubedex, "Prometheus-Rabbitmq-Exporter," [Online]. Available: <https://kubedex.com/resource/prometheus-rabbitmq-exporter/>.
- [15] "Github Docs of Postgres exporter," [Online]. Available: [https://docs.gitlab.com/ee/administration/monitoring/prometheus/postgres\\_exporter.html](https://docs.gitlab.com/ee/administration/monitoring/prometheus/postgres_exporter.html).
- [16] P. org, "About PostgreSQL," [Online]. Available: <https://www.postgresql.org/about/>.